

MACHINE LEARNING

COURSE OUTLINE

JS

August, 2022

Course title : Machine learning
Lecturer : Jonas Striaukas. Email: js.fi@cbs.dk or jonas.striaukas@gmail.com
Lecture time : TBA
Auditorium : TBA
Course website : https://jstriaukas.github.io/ml_course ↗

Important : please write *Machine learning class* in the subject line when you write to me regarding the course material/questions/etc. Please write questions after each lecture, I will allocate 10-15 minutes at the beginning of the following lecture to answer your questions.

Prerequisites : introduction to statistics, linear regression and some basic computing in statistical software (e.g., R, Python) is assumed, but otherwise it is a self-contained course. I recommend to read through the introductory book (first book in the list of recommended books) before the course to have a rough idea on what we will cover during the course.

Books: I try to make material self-contained so that you don't need to buy any book for the course. However, if you like the course and want to have a deeper understanding on a particular topic or machine learning in general, I suggest the following books:

- (great introductory book) James, G., Witten, D., Hastie, T., & Tibshirani, R. (2013). *An introduction to statistical learning* (Vol. 112, p. 18). Springer (New York).
 ► online copy: [pdf](#) ↗.
- (moderate level overview book) Hastie, T., & Friedman, J. H., Tibshirani, R. (2009). *The elements of statistical learning: data mining, inference, and prediction* (Vol. 2, pp. 1-758). Springer (New York).
- (covers more recent topics) Fan, J., Li, R., Zhang, C. H., & Zou, H. (2020). *Statistical foundations of data science*. Chapman and Hall/CRC.

Software: students need to install statistical software R or some other program (Python, Matlab, Julia, Octave, ...). I will be showing examples mainly in R and Python. I advise you to install the statistical software before the course starts (if you don't have it already). See below for the instructions how to install R and Python. You are free to use Matlab or Julia or other language of your choice to code the examples we cover in the class.

Exam: TBA

Information about the course

Topics

The course will cover the following main topics:

- **Topic 1:** Introduction to learning, multiple and nonparametric regression
▶ Material: [slides shinyapps](#) [slides pdf](#) [tablet friendly slides pdf](#)
- **Topic 2:** High-dimensional linear regression
▶ Material: [slides shinyapps](#) [slides pdf](#) [tablet friendly slides pdf](#)
- **Topic 3:** High-dimensional regression properties and generalized linear models (GAMs)
▶ Material: [slides shinyapps](#) [slides pdf](#) [tablet friendly slides pdf](#)
- **Topic 4:** Prediction, loss functions and M-estimators
▶ Material: [slides shinyapps](#) [slides pdf](#) [tablet friendly slides pdf](#)
- **Topic 5:** Introduction to deep learning
▶ Material: [slides shinyapps](#) [slides pdf](#) [tablet friendly slides pdf](#)
- **Topic 6:** Introduction to causal machine learning
▶ Material: [slides shinyapps](#) [slides pdf](#) [tablet friendly slides pdf](#)

Details on each topic

Instructions to install R

The main R software is available at <https://cran.r-project.org>. Once on this website, you need to select an file to download for your operating system (OS), so if you work on Windows you need to download and install Window `.exe` file. Please install the most recent version of R. I also strongly advise to install RStudio (free version of it) from <https://rstudio.com>. You need to install R prior to RStudio.

Install R packages: suppose you need to install an R package called ‘`forecast`’. You should write in RStudio console:

```
install.packages("forecast")
```